




Art & Science to running an AI Use Case Discovery Workshop

Vish Rajagopalan



The AI Use case Discovery Challenge

The AI Specialists Dilemma

BACKGROUND

- **AE Says** : AI is “Strategic” for the Customer. We need to “wow them” (*Renewal Coming up :)*)
- **SE** : Already shown them “AI Studios” (*They are curious but not sure about the product*)

Expectations

- “We need to get them to get understand **“value” of the product**”

- How do I **“discover”** what customer wants to do ?
- What **“domain”** questions should i ask ?

- Which use case serves best
- Do they have “right” data
- How to align expectations



AI Discovery Runbook

Very Draft – Need your updates today

What It is :


- **A Structured framework** that allows you to identify assets you need to build on how to run Use case Discovery workshops/ meetings
- **A “Think Tool”** : Provides Ideas around Workshop structure , contents, and prioritization approach

What It is **NOT**:

- **A ChatGPT of sorts , that “solves” the use case discovery for you:**
 - You still need to work with right stakeholders
 - Do sniff tests
 - Do Whiteboarding and all the good stuff

The Anatomy of a Discovery Workshop

Playbook on a slide.



Pre Discovery Homework: Conduct Interview/Survey

Owner : Field Specialist. Use this when the Customer /SE is not sure what problem to tackle



IDENTIFY THE FRICTION

Focusing on the pain points where AI typically excels (Repetition, Data Overload, Poor Prediction).

- **The Identification:** Identify one core process in your department that consumes at least 20% of your team's weekly manual effort.
 - *Examples: Manually categorizing emails, extracting data from PDFs, summarizing long reports, or verifying invoices.*
- **The Bottleneck:** In the process mentioned above, what is the primary "blocker"?
 - ☐ Waiting for human approval/review.
 - ☐ Accessing data from multiple disparate systems.
 - ☐ Cleaning or formatting inconsistent data.
 - ☐ Complex decision-making that requires specialized knowledge.
- **The "If I Only Knew" Question:** If you could predict one business variable with 90% accuracy (e.g., customer churn, equipment failure, next month's demand), which one would change your strategy the most?

IMPORTANT : Use open ended questions, and use GPT tools to get deeper in the customer domain

Pre Discovery Homework: Conduct Interview/Survey

Owner : Field Specialist. Use this when the Customer /SE is not sure what problem to tackle



Data & Tooling Readiness

Assess availability of the data, controls and approvals needed

- **Data Format:** Where does the information for this process currently live? (Select all that apply)
 - ☐ Structured Databases (SQL, Snowflake, ERP)
 - ☐ Unstructured Documents (PDFs, Word, Slide decks)
 - ☐ Internal Communications (Slack, Emails, Teams)
 - ☐ Tribal Knowledge (Only lives in people's heads)
- **Data Quality:** On a scale of 1–5, how much do you trust the accuracy and cleanliness of this data today?
 - (1 = Total Mess, 5 = Ready for Analysis)

IMPORTANT : Data Quality and Readiness will be your #1 area where you could get quickly derailed

Pre Discovery Homework: Conduct Interview/Survey

Owner : Field Specialist. Use this when the Customer /SE is not sure what problem to tackle



Strategic Impact

Helps in prioritization across use cases

- **The Value Driver:** If we automated or "augmented" this task with AI, what would be the primary win?
 - ☐ **Cost Savings:** Reducing headcount or operational spend.
 - ☐ **Velocity:** Doing the same task 10x faster.
 - ☐ **Quality:** Reducing human error and increasing consistency.
 - ☐ **Revenue:** Creating a new product or better customer experience.
- **Risk Tolerance:** How "accurate" does this AI need to be for you to trust it?
 - ☐ **Low Risk:** 70% accuracy is fine (e.g., internal drafting).
 - ☐ **High Risk:** 99% accuracy is required (e.g., financial reporting, legal compliance)

IMPORTANT : Identifying the Strategic Impact helps the customer also internally sell the project

Pre Discovery Homework: Conduct Interview/Survey

Owner : Field Specialist. Use this when the Customer /SE is not sure what problem to tackle



Blue Sky Idea

Smart Question to ask in case the customer has an “hidden agenda” and want their own “project” to be prioritized

- **Blue Sky Thinking:** If you had a "Magic Assistant" that knew everything about your Department/Domains historical data and could answer any question in seconds, what is the first question you would ask it?

IMPORTANT : Identifying the Strategic Impact helps the customer also internally sell the project, **Also build your own questionnaire list based on domain or engage Industry Solutions team**

CLOUDERA



Workshop Participants

Getting the right people in, who matter

- **The Profile:**
 - **Executive Sponsor:** To ensure alignment with North Star goals.
 - **Product/ Business Owners (SMEs):** The "problem owners" who know the day-to-day pain.
 - **IT/Data Architects:** To check data availability and security "red flags" early.
 - **End Users:** To ensure the solution is actually adoptable.
- **Interview/Survey Strategy:**
 - **Pre-work:** Survey Questionnaire or Interviews
 - **Tools:** Use tools like Miro or white boards
 - **"Kill the Idea":** Ask interviewers to explain reasons why a use case won't work to build robustness.



Workshop Mechanics

*Remember you are the “facilitator”,
let the customer own the show*

1. **Phase 1: Why are we here ? (45 min):** Workshop sponsor Identifies priorities, defines expectations,
2. **Phase 2: Problem Mining (60 min):** Identifying friction points in the current workflow for the use cases identified in the survey/interview (not starting with the tech).
3. **Phase 3: Prioritization (45 min):** Plotting ideas on the Value vs. Feasibility matrix.
4. **Phase 4: Solution Ideation (60 min) / Picked Use Case:** Mapping AI capabilities to those friction points using the AI Canvas and build first draft Reference Architecture.
5. **Phase 5: Roadmap & Ownership (30 min):** Defining the MVP and next steps.



Fleshing out the Use case

Think of atleast 4 categories of questions per use case


1. **Business Impact & Value :**
 - How frequently does this problem occur (daily, weekly, ad hoc)?
 - What is the current cost of not solving the problem
2. **Process & Workflow Understanding**
 - Where are the **biggest delays or bottlenecks** today?
 - Which steps are repetitive vs. judgment-based?
3. **Data Readiness**
 - How far back does historical data go? Where do we stop?
 - Are there labels or outcomes available (e.g., success/failure, resolved/unresolved)?
4. **Risk and Compliance**
 - Are there regulatory, compliance, or audit constraints?
 - What happens if the AI system is unavailable?

Use structured Frameworks to prioritize across use cases



Prioritizing across use cases

Engage the participants to prioritize by using sticky notes across



AI Use case prioritization rubric or just ask !! :)



But How to Prioritize?

Use a scoring rubric or ask the participants directly

Category 1: Business Value (The "Why")

Criterion	1 Point (Low)	3 Points (Medium)	5 Points (High)
Strategic Alignment	"Nice to have" task.	Aligned with departmental goals.	Directly impacts "North Star" KPIs/CEO goals.
Frequency/Volume	Happens once a month.	Daily task for a small team.	High-volume, continuous process (Enterprise-wide).
Potential ROI	Marginal time savings.	Measurable cost reduction or 20%+ time back.	Revenue generation or massive risk mitigation.
User Experience	Internal-only, invisible.	Improves internal morale/speed.	Directly improves Customer Satisfaction (NPS/CSAT).

Category 2: Technical Feasibility (The "How")

Criterion	1 Point (Low)	3 Points (Medium)	5 Points (High)
Data Availability	Data is "tribal" or in paper files.	Data is in PDFs/Digital docs (RAG ready).	Data is structured in a clean API or Database.
Tolerance for Error	Needs 100% accuracy (e.g., Legal).	Human-in-the-loop required (e.g., Drafting).	Low risk if "hallucinations" occur (e.g., Creative).
Complexity of Logic	Requires deep "intuition" or feeling.	Follows clear business rules.	Pattern recognition or summarization tasks.
Integration Ease	Requires a total system overhaul.	Can be a standalone "side-car" app.	Plugs directly into existing tools (Slack/ERP).

Use the AI Canvas to dig deeper per prioritized use case



An Example with ARD (Australia team)
facilitated by Vish

AI Canvas for each Use case

Leave with all boxes filled

The AI Canvas: Intelligent Invoice Processing

What task/decision are you examining?

Automatically extract key structured attribute data (Invoice Number, Due Date, Supplier) from unstructured or semi-structured documents (PDFs, Scans) across multiple languages and varied formats.



Prediction

Identify the key uncertainty that you would like to resolve.

Is the extracted value for a specific attributes (e.g., Invoice Number) correct, and what is the confidence level of this extractions?



Judgment

Determine the payoffs to being right versus being wrong. Consider both false positives and false negatives.

False Positive (FP): The AI extracts the wrong value but assigns High Confidence. **Payoff:** Incorrect data enters ERP, requiring costly manual reconciliation/correction later. **False Negative (FN):** The AI extracts the correct value but assigns Low Confidence. **Payoff:** Document is flagged for manual review (HITL), wasting staff time but ensuring data quality. **Priority:** Minimize FPs, as FNs are caught by the HITL process.



Action

What are the actions that can be chosen?

- Invoice Metadata with high confidence of extraction above a threshold should have Straight through Processing o ERP.
- Invoice Data with low confidence should be in an Exception Management Queue for Human Processing
- Human Processed Invoices should be used for Retraining



Outcome

Choose the measure of performance that you want to use to judge whether you are achieving your outcomes.

1. **F1 Score/Accuracy:** Must be >95% on production data. 2. **Straight-Through Processing (STP) Rate:** Percentage of documents processed without human intervention. 3. **Time-to-Value:** Average time from document ingestion to data available for retrieval.



Training

What data do you need on past inputs, actions and outcomes in order to train your AI and generate better predictions?

Historical PDFs/Scans, along with the corresponding correct attribute values (labeled data). The key is the user-corrected data from the HITL process.



Input

What data do you need to generate predictions once you have an AI algorithm trained?

Raw PDF or image file (Ingested via Email, or UI drop).



Feedback

How can you use measured outcomes along with input data to generate improvements to your predictive algorithm?

The manual corrections made by AP staff in the HITL Exception UI are captured, labeled, and used to continuously retrain and fine-tune the AI model on new or problematic document formats.

How will this AI impact on the overall workflow?

Explain here how the AI for this task/decision will impact on related tasks in the overall workflow Will it involve staff retraining or job redesign?

The role of the AP/Sales staff shifts from **Data Entry** to **Exception Handling, Validation, and Model Training**. This requires job redesign and retraining for efficiency gains.


Use the AI Canvas to dig deeper per prioritized use case



An Example with ARD (Australia team)
facilitated by Vish)

Use tools to get to a first draft architecture

Leave with all boxes filled






Workshop Outcomes and Next Steps



Aim for
tangible
outcomes
and Buy Ins

*Avoid concluding without clear next
steps on what happens next*



1. **A Prioritized Backlog:** A list of 3-5 "Ready-to-Build" use cases.
2. **High-Level Solution Design:** see example
3. **Identify what happens next :** Build a mental RACI Matrix on who owns what going forward




RACI - Matrix

	Person A	Person B	Person C	Person D	
Task 1	R				Responsible
Task 2		A			Accountable
Task 3			C		Consulted
Task 4				I	Informed

A GOOD READING



<https://cdn.openai.com/business-guides-and-resources/identifying-and-scaling-ai-use-cases.pdf>



Backup

CLOUDERA

CLOUDERA